# Normal Approximations (part I)

Ryan Miller

Normal Distributions

Having now seen many distributions, you may have noticed the prevalence of a certain shape:

1000 bootstrap samples

- ▶ Nearly all of the bootstrap distributions we've seen are symmetric and bell-shaped
- ▶ This is not a coincidence, it's backed up by statistical theory

Normal Distributions

▶ These distributions can be characterized by the curve:

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

▶ This curve defines the **Normal Distribution**
  - ▶ μ is the center (mean) of the distribution
  - ▶ σ is the standard deviation of the distribution
  - ▶ We use the shorthand $N(\mu, \sigma)$ to express a normal distribution, for example: N(3,1) is a curve centered at 3 with a standard deviation of 1

▶ You don't need to know the formula for the normal curve, but you should know that it depends on μ and σ

Normal Approximation

▶ As an alternative to bootstrapping, we can use the normal curve to approximate the *sampling distribution*

▶ To construct this approximation, we need to determine the mean and standard deviation of appropriate normal curve (any normal curve is entirely defined by μ and σ)

▶ Based upon what we know about sampling distributions, the approximation should look like $N(\text{estimate}, \text{SE})$

    ▶ Determining the *SE* leaves us in the same predicament that forced us to turn to bootstrapping

The Central Limit Theorem

We won't get into the details, but the **Central Limit Theorem** (CLT), which is one of the most well-known results in the history of statistics, provides a mathematical expression for the *SE* of many common statistics!

We'll first look at a CLT result for one proportion:

$$\hat{p} \sim N\left(p, \sqrt{\frac{p(1-p)}{n}}\right)$$

In words, the sample proportion $\hat{p}$ follows a normal distribution with a mean of p and standard deviation of $\sqrt{\frac{p(1-p)}{n}}$, thus providing a normal approximation of the sampling distribution

Using the CLT (one proportion)

$$\hat{p} \sim N\left(p, \sqrt{\frac{p(1-p)}{n}}\right)$$

- This suggests $SE = \sqrt{\frac{p(1-p)}{n}}$ when estimating a single proportion
- While we don't know p, $\hat{p}$ is our *best estimate*
- Which leads to the 95% confidence interval:

$$\hat{p} \pm 2 * \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

Confidence Interval Coverage

Let's check whether these 95% confidence intervals actually work by using random samples of size 20 to estimate the proportion of Fall 2018 Sta-209 students who "took the class for fun".

| Sample ID	Sample proportion	Calculation	95% CI
1	0.35	0.35 +/- 2* 0.107	(0.137,0.563)
2	0.25	0.25 +/- 2* 0.097	(0.056,0.444)
3	0.4	0.4 +/- 2* 0.11	(0.181,0.619)
4	0.55	0.55 +/- 2* 0.111	(0.328,0.772)
5	0.35	0.35 +/- 2* 0.107	(0.137,0.563)
6	0.25	0.25 +/- 2* 0.097	(0.056,0.444)

Confidence Interval Coverage

Of the first 200 confidence intervals, 3 fail to capture the true p,
suggesting the procedure is valid (and perhaps slightly conservative)

Confidence Interval Coverage

This finding should make sense, the normal curve does a good job approximating the actual sampling distribution in this example:

Sampling Distribution of $\hat{p}$

Practice

In a study conducted by Johns Hopkins University researchers investigated the survival of babies born prematurely. They searched their hospital's medical records and found 39 babies born at 25 weeks gestation (15 weeks early), 31 of these babies went on to survive at least 6 months. With your group:

1. Use a normal approximation to construct a 95% confidence interval estimate for the true proportion of babies born at 25 gestation that are expected to survive.
2. An article on Wikipedia suggests 70% of babies born at 25 weeks gestation survive. Is this claim consistent with the Johns Hopkins study?

Practice - Solution

1. $\hat{p} = 31/39 = 0.795$, using the normal approximation provided by CLT, $SE = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}} = \sqrt{\frac{0.795(1-0.795)}{39}} = 0.065$; this suggests the 95% CI:

$$0.795 \pm 2 * 0.065 = (0.668, 0.922)$$

2. Yes, 0.70 is contained in the 95% confidence interval, suggesting it is a plausible value of the population parameter.

Sufficiently Large?

The normal approximation suggested by the Central Limit Theorem only works well when n is sufficiently large

▶ For a single proportion, "sufficiently large" also depends upon the value of p
▶ A common rule of thumb for whether the normal approximation of $\hat{p}$ is reasonable requires:

1. $n * p \geq 10$
2. $n * (1 - p) \geq 10$

If *either* of these conditions isn't met, you should consider another approach (such as bootstrapping)

Using the CLT (difference in proportions)

For a difference in proportions, provided n_1 and n_2 are sufficiently large, CLT suggests:

$$\hat{p}_1 - \hat{p}_2 \sim N\left(p_1 - p_2, \sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}\right)$$

In this scenario, checking the conditions of this normal approximation is a little more tedious:

1. $n_1 * p_1 \geq 10$
2. $n_1 * (1 - p_1) \geq 10$
3. $n_2 * p_2 \geq 10$
4. $n_2 * (1 - p_2) \geq 10$

Practice

A common stereotype is that medical doctors have terrible handwriting. If true, this might contribute to errors in medical prescriptions that rely on hand-written forms. A 2010 study took two groups of doctors that had similar error rates before the study and randomly assigned half of them to use an electronic prescription form, while the other half continued using written prescriptions. After 1 year, the error rate of each group was recorded:

	Error	Non-errors	Total
Electronic	254	3594	3848
Hand-written	1478	2370	3848

1. Find the 95% confidence interval for difference proportions (e-prescriptions minus hand-written prescriptions)
2. Is it plausible that both electronic and handwritten forms have the same error rate?

Practice - Solution

1. $\hat{p}_e - \hat{p}_{hw} = 254/2848 - 1478/3848 = 0.066 - 0.384 = -0.318$;
   while $SE = \sqrt{\frac{0.066(1-0.066)}{3848} + \frac{0.384(1-0.384)}{3848}} = 0.017$; thus the
   95% Confidence Interval is given by:

$$-0.318 \pm 2 * 0.017 = (-0.352, -0.284)$$

2. No, the 95% confidence interval does not contain zero,
   implying the error rates being equal for both forms is not
   plausible

Confidence Levels that aren't 95%

Recall that confidence intervals have the form:

$$\text{Estimate} \pm c * SE$$

▶ Because areas under the normal curve are known, we aren't limited by the 68-95-99 rule when it comes to determining a meaningful value for c

▶ The **standard normal** distribution has a mean of 0 and standard deviation of 1

    ▶ We can use cut-points in this distribution to achieve *any* confidence level that we'd like

    ▶ One place we can do this is the "Theoretical Distribution" menu on StatKey

Practice

Recall that in the handwritten vs. electronic prescription study:
$\hat{p}_e - \hat{p}_{hw} = 254/2848 - 1478/3848 = 0.066 - 0.384 = -0.318$ and
$SE = \sqrt{\frac{0.066(1-0.066)}{3848} + \frac{0.384(1-0.384)}{3848}} = 0.017$

1. Use StatKey to find the appropriate normal quantile (ie: c) for constructing an 80% confidence interval
2. Use this value to construct an 80% confidence interval for effect of handwritten vs. electronic prescriptions

Practice - Solution

1. $c = 1.282$
2. $-0.318 \pm 1.282 * 0.017 = (-0.340, -0.296)$

Conclusion

Right now you should. . .

1. Understand how a normal approximation can be used to describe the sampling distributions
2. Know how to use a normal approximation to construct confidence intervals for scenarios involving one proportion, or a difference in proportion
3. Be aware of the assumptions required for the normal approximation be reasonable

If you want more information:

▶ Read Ch 5.2, Ch 6.1, and Ch 6.3